

Deep Learning Homework 5 Report

Cheng-Liang Chi

April 30, 2025

Contents

1	Introduction	3
2	Implementation Details	3
2.1	Refactoring the Sample Code	3
2.1.1	DQN Class	3
2.1.2	DQNAgent Class	4
2.1.3	Trainer Class	7
2.1.4	Tester Class	11
2.2	Hyperparameters	13
2.3	Bellman Equation	13
2.4	Prioritized Experience Replay	15
2.5	Multi-Step Reward	15
2.6	Weight & Biases Using Techniques	17
2.6.1	Categorizing the Results	17
2.6.2	Snapshotting the Code	17
2.6.3	Hyperparameter Tuning	17
2.6.4	Result Plot Generation	17
2.6.5	WandB Initialization	18
3	Discussion	18
3.1	Task 1: CartPole-v1 with vanilla DQN	18
3.1.1	Training Commands	18
3.1.2	Training Curves	19
3.1.3	Testing Commands	20
3.2	Task 2: ALE/Pong-v5 with vanilla DQN	20
3.2.1	Training Commands	20
3.2.2	Training Curves	21
3.2.3	Testing Commands	21
3.3	Task 3: ALE/Pong-v5 with DQN Variants	21
3.3.1	Training Commands	21
3.3.2	Training Curves	22
3.3.3	Testing Commands	23
3.4	Analyze each Technique	23
3.4.1	Double DQN	23
3.5	Prioritized Experience Replay	24

3.6	Multi-Step Learning	24
3.7	Additional Training Tricks	24
3.7.1	Multiple Evaluation Episodes	24
3.7.2	Learning Rate Scheduler	24

1 Introduction

In this report, we will discuss the implementation of Deep Q-Learning (DQN) and its variants, including Double DQN [1], Prioritized Experience Replay [2], and Multi-Step Learning [3]. We will also analyze the performance of these algorithms on both the `CartPole-v1` and `ALE/Pong-v5` environments. The goal is to understand the impact of these techniques on the learning process and the final performance of the agent.

2 Implementation Details

2.1 Refactoring the Sample Code

First, since the sample code is not modular, we refactored the code into multiple classes. The main classes are:

- **DQN:** This class implements the neural network architecture. It is a subclass of the `torch.nn.Module` class.
- **DQNAgent:** This class represents the agent that interacts with the environment. It contains methods for selecting actions, storing experiences, and updating the Q-values. The agent is responsible for the exploration and exploitation of the environment. It also contains the experience replay buffer, which stores the agent's experiences for training. However, the backend neural network is modularized into a separate class. That is we can reuse the same agent class for different neural networks dealing with different environments we focused on.
- **Trainer:** This class is responsible for training the agent. It handles the training loop, including collecting experiences, updating the model, and logging results. By the way, the trainer class is also responsible for loading the environment and the agent. It will handle agent's actions and interact it with the environment.
- **Tester:** This class is responsible for testing the agent. It handles the testing loop, it is similar to the trainer class, but it does not update the model. The tester class is also responsible for loading the environment and the agent. It will handle agent's actions and interact it with the environment.

With this refactoring, I can easily add new features and techniques to the agent, trainer, and tester classes without modifying the core logic of the DQN algorithm. This modularity also allows us to reuse the same code for different tasks and environments, making it easier to experiment with different architectures and techniques.

2.1.1 DQN Class

The DQN class is a subclass of the `torch.nn.Module` class. It implements the neural network architecture for the DQN algorithm.

For the `CartPole-v1` environment, I used a simple feedforward neural network with two hidden layers. The input layer has 4 neurons (one for each state variable), and the output layer has 2 neurons (one for each action). Both hidden layers have 128 neurons, and the activation function is ReLU. The architecture is designed to be simple and efficient, as the `CartPole-v1` environment has a relatively low-dimensional state space. The output layer uses

a linear activation function, which is suitable for the DQN algorithm since it outputs Q-values for each action.

The implementation of the CartPoleDQN class is as follows:

```
5 class CartPoleDQN(nn.Module):
6     def __init__(self, input_state, num_actions):
7         super().__init__()
8         self.network = nn.Sequential(
9             nn.Linear(input_state, 128),
10            nn.ReLU(),
11            nn.Linear(128, 128),
12            nn.ReLU(),
13            nn.Linear(128, num_actions),
14        )
15
16    def forward(self, x):
17        return self.network(x)
```

For the ALE/Pong-v5 environment, I used a convolutional neural network (CNN) with three convolutional layers and two fully connected layers. The architecture is designed to effectively capture the spatial features of the input frames, allowing the agent to learn optimal policies for playing the game.

The implementation of the PongDQN class is as follows:

```
20 class PongDQN(nn.Module):
21     def __init__(self, input_channels, num_actions):
22         super().__init__()
23         self.network = nn.Sequential(
24             nn.Conv2d(input_channels, 32, kernel_size=8, stride=4),
25             nn.ReLU(),
26             nn.Conv2d(32, 64, kernel_size=4, stride=2),
27             nn.ReLU(),
28             nn.Conv2d(64, 64, kernel_size=3, stride=1),
29             nn.ReLU(),
30             nn.Flatten(),
31             nn.Linear(64 * 7 * 7, 512),
32             nn.ReLU(),
33             nn.Linear(512, num_actions),
34         )
35
36    def forward(self, x: torch.Tensor):
37        x = x.squeeze(-1)
38        return self.network(x / 255.0)
```

2.1.2 DQNAgent Class

The DQNAgent class is responsible for interacting with the environment and managing the agent's experiences. It contains methods for selecting actions, storing experiences, and updating

the Q-values.

The agent uses an ϵ -greedy policy for action selection, where it explores the environment with a probability of ϵ and exploits the learned Q-values with a probability of $1 - \epsilon$. The agent also maintains a replay buffer, which stores the agent's experiences for training. The replay buffer is implemented using the `PrioritizedReplayBuffer` class, which will be discussed in Section 2.4.

The implementation of the `DQNAgent` class is as follows:

```
85 class DQNAgent:
86     def __init__(self, env, args):
87         self.device = args.device
88         logger.info(f"Using device: {self.device}")
89
90         match env.spec.id:
91             case "CartPole-v1":
92                 self.input_state = 4
93                 self.num_actions = 2
94                 self.DQN = CartPoleDQN
95             case "ALE/Pong-v5":
96                 self.input_state = 4
97                 self.num_actions = 6
98                 self.DQN = PongDQN
99             case _:
100                 raise ValueError(f"Unsupported environment: {env}")
101
102         self.q_net = self.DQN(self.input_state,
103                               ↪ self.num_actions).to(self.device)
104         self.q_net.apply(init_weights)
105
106         self.save_dir = args.save_dir
107         os.makedirs(self.save_dir, exist_ok=True)
108
109     def train(self, args):
110         self.batch_size = args.batch_size
111         self.update_period = args.update_period
112         self.gamma = args.discount_factor
113         self.memory = PrioritizedReplayBuffer(
114             args.memory_size, args.per_alpha, args.per_beta
115         )
116
117         self.vanilla = args.vanilla
118         self.target_net = self.DQN(self.input_state,
119                                     ↪ self.num_actions).to(self.device)
120         self.target_net.load_state_dict(self.q_net.state_dict())
121         self.target_net.eval()
122         self.optimizer = optim.Adam(self.q_net.parameters(), lr=args.lr)
123
124         self.learn_count = 0
125
126     def select_action(self, state, epsilon):
```

```

125     if random.random() < epsilon:
126         return random.randint(0, self.num_actions - 1)
127     state_tensor = (
128         ↪ torch.from_numpy(np.array(state)).float().unsqueeze(0).to(self.device)
129     )
130     self.q_net.eval()
131     with torch.no_grad():
132         q_values = self.q_net(state_tensor)
133     self.q_net.train()
134     return q_values.argmax().item()
135
136 def learn(self) -> float:
137     if len(self.memory) < self.batch_size:
138         return float("-inf")
139
140     batch, indices, weights = self.memory.sample(self.batch_size)
141     indices = torch.tensor(indices, dtype=torch.int64).to(self.device)
142     weights = torch.tensor(weights,
143         ↪ dtype=torch.float32).to(self.device)
144     states, actions, rewards, next_states, dones = zip(*batch)
145
146     states =
147         ↪ torch.from_numpy(np.array(states).astype(np.float32)).to(self.device)
148     next_states =
149         ↪ torch.from_numpy(np.array(next_states).astype(np.float32)).to(
150             self.device
151         )
152     actions = torch.tensor(actions, dtype=torch.int64).to(self.device)
153     rewards = torch.tensor(rewards,
154         ↪ dtype=torch.float32).to(self.device)
155     dones = torch.tensor(dones, dtype=torch.float32).to(self.device)
156
157     with torch.no_grad():
158         if self.vanilla:
159             target_q_values = (
160                 rewards
161                 + (1 - dones) * self.gamma *
162                 ↪ self.target_net(next_states).max(1)[0]
163             )
164         else:
165             next_q_values = self.target_net(next_states).max(1)[0]
166             target_q_values = rewards + (1 - dones) * self.gamma *
167             ↪ next_q_values
168     q_values = self.q_net(states).gather(1,
169         ↪ actions.unsqueeze(1)).squeeze(1)
170
171     td_errors = target_q_values - q_values
172     loss = (td_errors**2 * weights).mean()
173     # loss = nn.MSELoss()(q_values, target_q_values)

```

```

167     self.optimizer.zero_grad()
168     loss.backward()
169     # torch.nn.utils.clip_grad_norm_(self.q_net.parameters(), 1.0)
170     self.optimizer.step()
171
172     self.memory.update_priorities(indices,
173     ↪     td_errors.detach().cpu().numpy())
174
175     # self.memory.beta = min(1.0, self.memory.beta + 0.000001)
176
177     self.learn_count += 1
178     if self.learn_count % self.update_period == 0:
179         self.target_net.load_state_dict(self.q_net.state_dict())
180         logger.debug(
181             f"Target network updated at
182             ↪ learn_count={self.learn_count/1000:.2f}k"
183         )
184         logger.debug(
185             f"Memory size: {len(self.memory)}, beta:
186             ↪ {self.memory.beta:.2f}"
187         )
188     return loss.item()

```

2.1.3 Trainer Class

The Trainer class is responsible for training the agent. It handles the training loop, including collecting experiences, updating the model, and logging results. The trainer class is also responsible for loading the environment and the agent. It will handle agent's actions and interact it with the environment. The implementation of the Trainer class is as follows:

```

18 class Trainer:
19     def __init__(self, args) -> None:
20         self.save_dir = args.save_dir
21         self.env = gym.make(args.env_name, render_mode="rgb_array")
22         if args.env == "pong":
23             self.env = AtariPreprocessing(
24                 self.env,
25                 frame_skip=1,
26                 grayscale_newaxis=True,
27                 screen_size=84,
28                 grayscale_obs=True,
29                 noop_max=30,
30             )
31             self.env = FrameStackObservation(self.env, 4)
32         self.env.action_space.seed(args.seed)
33         self.env.observation_space.seed(args.seed)
34         self.env.reset(seed=args.seed)
35
36         self.num_actions = self.env.action_space.n # type: ignore

```

```

37     logger.info(f"Environment: {self.env.spec.id}") # type: ignore
38     logger.info(f"Action Space: {self.env.action_space}")
39     logger.info(f"Observation Space: {self.env.observation_space}")
40
41     self.device = torch.device("cuda" if torch.cuda.is_available() else
42                               ↪ "cpu")
43     logger.info(f"Using device: {self.device}")
44
45     self.agent = DQNAgent(self.env, args=args)
46     self.agent.train(args)
47     self.preprocessor = DummyPreprocessor()
48
49     self.epsilon = args.epsilon_start
50     self.epsilon_decay = args.epsilon_decay
51     self.epsilon_min = args.epsilon_min
52
53     self.episode = 0
54     self.env_step = 0
55     self.best_reward = 0 if self.env == "cartpole" else -21
56
57     self.learn_per_step = args.learn_per_step
58
59     self.eval_episodes = args.eval_episodes
60
61     def run(self, episodes=1000):
62         with Progress(
63             SpinnerColumn(),
64             *Progress.get_default_columns(),
65             TimeElapsedColumn(),
66             MofNCompleteColumn(),
67             console=console,
68         ) as progress:
69             task = progress.add_task("[cyan]Training...", total=episodes)
70             eval_task = progress.add_task(
71                 "[cyan]Episode 0: Evaluating...",
72                 total=self.eval_episodes,
73             )
74             for ep in range(episodes):
75                 progress.update(task, description=f"[cyan]Episode {ep}:
76                 ↪ Training...")
77                 self.episode = ep
78                 if self.env_step > 20e6:
79                     logger.info(f"Reached 20M steps, stopping training.")
80                     break
81                 if (ep + 1) % 20 == 0:
82                     logger.info(
83                         f"Episode {self.episode}: Environment Step:
84                         ↪ {self.env_step/1000:.2f}k, epsilon:
85                         ↪ {self.epsilon:.4f}"
86                     )

```



```

83
84     self.train()
85
86     if (ep + 1) % (episodes // 20) == 0:
87         model_path = os.path.join(self.save_dir,
88             ↪ f"model_ep{ep}.pt")
89         torch.save(self.agent.q_net.state_dict(), model_path)
90         logger.info(f"Saved model checkpoint to {model_path}")
91
92     if (ep + 1) % (episodes // 50) == 0:
93         eval_rewards = []
94         progress.reset(eval_task)
95         progress.update(
96             eval_task, description=f"[cyan]Episode {ep}:
97             ↪ Evaluating..."
98         )
99         for _ in range(self.eval_episodes):
100             eval_rewards.append(self.evaluate())
101             progress.update(eval_task, advance=1)
102         eval_reward = sum(eval_rewards) / len(eval_rewards)
103         max_eval_reward = max(eval_rewards)
104         logger.info(
105             f"Episode {ep} - Eval Reward: {eval_reward:.2f}
106             ↪ {eval_rewards}"
107         )
108
109     if eval_reward > self.best_reward:
110         self.best_reward = eval_reward
111         model_path = os.path.join(self.save_dir,
112             ↪ "best_model.pt")
113         torch.save(self.agent.q_net.state_dict(),
114             ↪ model_path)
115         logger.info(
116             f"Saved new best model to {model_path} with
117             ↪ reward {eval_reward}"
118         )
119     wandb.log(
120         {
121             "Episode": ep,
122             "Env Step Count": self.env_step,
123             "Eval Reward": eval_reward,
124             "Max Eval Reward": max_eval_reward,
125         }
126     )
127     progress.update(task, advance=1)
128
129 def train(self):
130     avg_loss = 0
131
132     obs, _ = self.env.reset(seed=random.randint(0, 10000))

```

```

127     state = self.preprocessor.reset(obs)
128
129     done = False
130     total_reward = 0
131
132     while not done:
133         action = self.agent.select_action(state, self.epsilon)
134         next_obs, reward, terminated, truncated, _ =
            ↪ self.env.step(action)
135         next_state = self.preprocessor.step(next_obs)
136         done = terminated or truncated
137
138         self.agent.memory.add((state, action, reward, next_state,
            ↪ done), 1)
139         state = next_state
140         total_reward += float(reward)
141         self.env_step += 1
142
143         for _ in range(self.learn_per_step):
144             loss = self.agent.learn()
145             if loss != float("-inf") and self.epsilon >
                ↪ self.epsilon_min:
146                 self.epsilon *= self.epsilon_decay
147             if loss != float("-inf"):
148                 avg_loss += loss
149         avg_loss /= self.learn_per_step
150         wandb.log(
151             {
152                 "Episode": self.episode,
153                 "Env Step Count": self.env_step,
154                 "Total Reward": total_reward,
155                 "Epsilon": self.epsilon,
156                 "Loss": avg_loss,
157             }
158         )
159
160     def evaluate(self):
161         obs, _ = self.env.reset(seed=random.randint(0, 10000))
162         state = self.preprocessor.reset(obs)
163         frames = [self.env.render()]
164
165         done = False
166         total_reward = 0
167
168         while not done:
169             action = self.agent.select_action(state, 0.0)
170             next_obs, reward, terminated, truncated, _ =
                ↪ self.env.step(action)
171             state = self.preprocessor.step(next_obs)
172             done = terminated or truncated

```

```

173         total_reward += float(reward)
174         frames.append(self.env.render())
175
176     return total_reward

```

2.1.4 Tester Class

The Tester class is responsible for testing the agent. It handles the testing loop, it is similar to the trainer class, but it does not update the model. The tester class is also responsible for loading the environment and the agent. It will handle agent's actions and interact it with the environment. The implementation of the Tester class is as follows:

```

19 class Tester:
20     def __init__(self, args) -> None:
21         self.save_dir = args.save_dir
22         self.env = gym.make(args.env_name, render_mode="rgb_array")
23         if args.env == "pong":
24             self.env = AtariPreprocessing(
25                 self.env,
26                 frame_skip=1,
27                 grayscale_newaxis=True,
28                 screen_size=84,
29                 grayscale_obs=True,
30                 noop_max=30,
31             )
32             self.env = FrameStackObservation(self.env, 4)
33
34         self.num_actions = self.env.action_space.n # type: ignore
35
36         self.device = torch.device("cuda" if torch.cuda.is_available() else
37                                     ↪ "cpu")
38         logger.info(f"Using device: {self.device}")
39         self.agent = DQNAgent(self.env, args=args)
40         self.agent.q_net.load_state_dict(
41             torch.load(args.model_path, map_location=self.device)
42         )
43         self.agent.q_net.to(self.device)
44         self.agent.q_net.eval()
45
46         self.preprocessor = DummyPreprocessor()
47
48         self.episode = 0
49         self.best_reward = 0 if self.env == "cartpole" else -21
50
51         self.visualize = args.visualize
52         self.save_dir = args.save_dir
53         self.seed = args.seed
54         gif_dir = os.path.join(self.save_dir, "gifs")

```

```

55     os.makedirs(gif_dir, exist_ok=True)
56
57     def run(self, episodes):
58         with Progress(
59             SpinnerColumn(),
60             *Progress.get_default_columns(),
61             TimeElapsedColumn(),
62             MofNCompleteColumn(),
63             console=console,
64         ) as progress:
65             task = progress.add_task("[cyan]Testing...", total=episodes)
66             total_reward = 0
67             for ep in range(episodes):
68                 self.episode = ep
69                 eval_reward = self.evaluate(seed=self.seed + ep)
70                 logger.info(f"Episode {ep} - Test Reward:
71                     ↳ {eval_reward:.2f}")
72                 total_reward += eval_reward
73                 wandb.log(
74                     {
75                         "Episode": ep,
76                         "Test Reward": eval_reward,
77                     }
78                 )
79                 progress.update(task, advance=1)
80             total_reward /= episodes
81             logger.info(f"Average Test Reward: {total_reward}")
82
83     def evaluate(self, seed):
84         obs, _ = self.env.reset(seed=seed)
85         state = self.preprocessor.reset(obs)
86         frames = [self.env.render()]
87
88         done = False
89         total_reward = 0
90
91         while not done:
92             action = self.agent.select_action(state, 0.0)
93             next_obs, reward, terminated, truncated, _ =
94                 ↳ self.env.step(action)
95             done = terminated or truncated
96             total_reward += float(reward)
97             state = self.preprocessor.step(next_obs)
98             frames.append(self.env.render())
99
100         if self.visualize:
101             gif_path = os.path.join(self.save_dir, "gifs",
102                 ↳ f"test_{self.episode}.gif")
103             imageio.mimsave(gif_path, frames, fps=30) # type: ignore
104             logger.info(f"Saved test episode frames to {gif_path}")

```

```
return total_reward
```

2.2 Hyperparameters

The hyperparameters used in the training process are crucial for the performance of the DQN algorithm.

For Tasks 1 and 2, we are required to use vanilla DQN without any modifications. To meet this requirement, certain hyperparameters must be set accordingly. For example, setting $\alpha = 0$ in the `PrioritizedReplayBuffer` makes the buffer behave like a uniform replay buffer. Additionally, setting `update_period = 1` in the `DQNAgent` class ensures that the model is updated every time a batch is sampled from the replay buffer, aligning with the standard behavior of vanilla DQN. And last, setting `n_steps=1` in the `DQNAgent` class ensures that the agent uses one-step returns for training, which is standard in vanilla DQN.

2.3 Bellman Equation

The Bellman equation is a fundamental concept in reinforcement learning that describes the relationship between the value of a state and the values of its successor states. In the context of DQN, the Bellman equation is used to update the Q-values based on the agent's experiences. The Bellman equation for Q-learning is given by:

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a') \quad (1)$$

where:

- $Q(s, a)$ is the Q-value for state s and action a .
- r is the reward received after taking action a in state s .
- γ is the discount factor, which determines the importance of future rewards.
- s' is the next state after taking action a in state s .
- $\max_{a'} Q(s', a')$ is the maximum Q-value for the next state s' over all possible actions a' .

The DQN algorithm uses a neural network to approximate the Q-values, and the Bellman equation is used to update the weights of the network during training. The Q-value update is performed using the following loss function:

$$L(\theta) = \mathbb{E}_{(s,a,r,s') \sim D} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2 \right] \quad (2)$$

where:

- $L(\theta)$ is the loss function.
- θ are the weights of the Q-network.
- θ^- are the weights of the target network, which are updated periodically.

- D is the replay buffer containing the agent's experiences.
- \mathbb{E} is the expectation operator, which averages over the experiences in the replay buffer.

The loss function measures the difference between the predicted Q-value and the target Q-value. The weights of the Q-network are updated using gradient descent to minimize this loss. The target Q-value is computed using the Bellman equation, and the target network is used to stabilize the training process. The target network is a separate copy of the Q-network, updated periodically every C steps (a hyperparameter) by copying the weights of the Q-network. This periodic update reduces the variance of the Q-value updates, making the training process more stable and efficient.

The implementation of the Bellman equation in the DQN algorithm is done in the `learn` method of the `DQNAgent` class. In this method, the Q-values are updated based on the agent's experiences, ensuring that the learning process aligns with the principles of reinforcement learning. The implementation of the `learn` method is as follows:

```

136 def learn(self) -> float:
137     if len(self.memory) < self.batch_size:
138         return float("-inf")
139
140     batch, indices, weights = self.memory.sample(self.batch_size)
141     indices = torch.tensor(indices, dtype=torch.int64).to(self.device)
142     weights = torch.tensor(weights,
143         ↪ dtype=torch.float32).to(self.device)
144     states, actions, rewards, next_states, dones = zip(*batch)
145
146     states =
147         ↪ torch.from_numpy(np.array(states).astype(np.float32)).to(self.device)
148     next_states =
149         ↪ torch.from_numpy(np.array(next_states).astype(np.float32)).to(
150             self.device
151         )
152     actions = torch.tensor(actions, dtype=torch.int64).to(self.device)
153     rewards = torch.tensor(rewards,
154         ↪ dtype=torch.float32).to(self.device)
155     dones = torch.tensor(dones, dtype=torch.float32).to(self.device)
156
157     with torch.no_grad():
158         if self.vanilla:
159             target_q_values = (
160                 rewards
161                 + (1 - dones) * self.gamma *
162                 ↪ self.target_net(next_states).max(1)[0]
163             )
164         else:
165             next_q_values = self.target_net(next_states).max(1)[0]
166             target_q_values = rewards + (1 - dones) * self.gamma *
167             ↪ next_q_values
168     q_values = self.q_net(states).gather(1,
169         ↪ actions.unsqueeze(1)).squeeze(1)

```

```

164     td_errors = target_q_values - q_values
165     loss = (td_errors**2 * weights).mean()
166     # loss = nn.MSELoss()(q_values, target_q_values)
167     self.optimizer.zero_grad()
168     loss.backward()
169     # torch.nn.utils.clip_grad_norm_(self.q_net.parameters(), 1.0)
170     self.optimizer.step()
171
172     self.memory.update_priorities(indices,
173     ↪ td_errors.detach().cpu().numpy())
174
175     # self.memory.beta = min(1.0, self.memory.beta + 0.000001)
176
177     self.learn_count += 1
178     if self.learn_count % self.update_period == 0:
179         self.target_net.load_state_dict(self.q_net.state_dict())
180         logger.debug(
181             f"Target network updated at
182             ↪ learn_count={self.learn_count/1000:.2f}k"
183         )
184         logger.debug(
185             f"Memory size: {len(self.memory)}, beta:
186             ↪ {self.memory.beta:.2f}"
187         )
188     return loss.item()

```

2.4 Prioritized Experience Replay

2.5 Multi-Step Reward

The multi-step reward is a technique used in reinforcement learning to improve the efficiency of learning by considering multiple steps of experience at once. In the context of DQN, multi-step rewards are used to update the Q-values based on a sequence of actions and rewards, rather than just the immediate reward. This approach allows the agent to learn from longer-term dependencies and can lead to faster convergence and better performance. The multi-step reward is computed by summing the rewards over a sequence of n steps, discounted by the discount factor γ :

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{n-1} r_{t+n-1} \quad (3)$$

where:

- R_t is the multi-step reward at time step t .
- r_t is the immediate reward at time step t .
- γ is the discount factor.
- n is the number of steps to consider for the multi-step reward.
- r_{t+k} is the immediate reward at time step $t + k$.

The multi-step reward is used to update the Q-values in the same way as the standard Bellman equation, but it incorporates the rewards from multiple steps. This allows the agent to learn from longer-term dependencies and can lead to faster convergence and better performance. The multi-step reward is implemented in the PrioritizedReplayBuffer class.

```

20 class PrioritizedReplayBuffer:
21     def __init__(self, capacity, alpha=0.6, beta=0.4, n_steps=3,
22         ↪ gamma=0.99):
23         self.capacity = capacity
24         self.alpha = alpha
25         self.beta = beta
26         self.n_steps = n_steps
27         self.gamma = gamma
28
29         self.buffer = []
30         self.priorities = np.zeros((capacity,), dtype=np.float32)
31         self.n_step_buffer = deque(maxlen=n_steps)
32         self.pos = 0
33
34     def add(self, transition, error):
35         self.n_step_buffer.append(transition)
36         if len(self.n_step_buffer) < self.n_steps:
37             return
38
39         # Compute multistep reward and next state
40         reward, next_state, done = self._get_multistep_transition()
41         state, action = self.n_step_buffer[0][:2]
42         multistep_transition = (state, action, reward, next_state, done)
43
44         priority = (abs(error) + 1e-6) ** self.alpha
45
46         if len(self.buffer) < self.capacity:
47             self.buffer.append(multistep_transition)
48         else:
49             self.buffer[self.pos] = multistep_transition
50
51             self.priorities[self.pos] = priority
52             self.pos = (self.pos + 1) % self.capacity
53
54     def _get_multistep_transition(self):
55         reward, next_state, done = 0, None, False
56         for idx, (_, _, r, ns, d) in enumerate(self.n_step_buffer):
57             reward += (self.gamma**idx) * r
58             next_state, done = ns, d
59             if done:
60                 break
61         return reward, next_state, done
62
63     def sample(self, batch_size):
64         assert len(self.buffer) >= batch_size, "Not enough samples to
65             ↪ sample from"

```



```

64
65     valid_size = len(self.buffer)
66     probs = self.priorities[:valid_size]
67     probs = probs / probs.sum()
68
69     indices = np.random.choice(valid_size, batch_size, p=probs)
70     samples = [self.buffer[i] for i in indices]
71
72     total = valid_size
73     weights = (total * probs[indices]) ** (-self.beta)
74     weights /= weights.max()
75     return samples, indices, weights.astype(np.float32)
76
77     def update_priorities(self, indices, errors):
78         for idx, err in zip(indices, errors):
79             self.priorities[idx] = (abs(err) + 1e-6) ** self.alpha
80
81     def __len__(self):
82         return len(self.buffer)

```

2.6 Weight & Biases Using Techniques

Since we are required to complete 3 different tasks in two different environments, I used the Weight & Biases (WandB) library to track the training process and visualize the results.

2.6.1 Categorizing the Results

To categorize the results, I used the `wandb.init` method to create a new run for each task and add tags to the run. This allows me to easily filter and compare the results of different runs in the WandB dashboard for the same task.

2.6.2 Snapshotting the Code

I also used set argument `save_code` to `True` and setting the `code_dir` argument to the whole current directory to save all `.py` files in the WandB run. This allows me to easily reproduce the results and compare the code used for different runs.

2.6.3 Hyperparameter Tuning

I used the `wandb.config` method to define the hyperparameters for each run. This allows me to easily track and compare the hyperparameters used for different runs in the WandB dashboard.

2.6.4 Result Plot Generation

I used the `wandb.log` method to log the results of each run. This allows me to easily visualize the results in the WandB dashboard and compare the performance of different runs. This also allows me to easily export the plots and use them in the report.

2.6.5 WandB Initialization

The implementation of the WandB initialization in the `Trainer` class is as follows:

```
241 wandb.init(  
242     project="DLP-Lab5-DQN",  
243     name=f"{args.exp}",  
244     tags=[args.env, "train"],  
245     save_code=True,  
246     settings=wandb.Settings(code_dir="."),  
247 )  
248 wandb.config.update(args)
```

And the implementation of the WandB initialization in the `Tester` class is as follows:

```
152     project="DLP-Lab5-DQN",  
153     name=f"{args.exp}",  
154     tags=[args.env, "test"],  
155     save_code=True,  
156     settings=wandb.Settings(code_dir="."),  
157 )  
158 wandb.config.update(args)
```

3 Discussion

3.1 Task 1: CartPole-v1 with vanilla DQN

3.1.1 Training Commands

For the task 1, which required to use the vanilla DQN algorithm, I used the following command to train the agent:

```
1 python3 trainer.py --env cartpole --exp report --vanilla --epsilon-decay  
↪ 0.99
```

The command specifies the environment as `CartPole-v1`, the experiment name as `report`, and the epsilon decay rate as 0.99. The `--vanilla` flag indicates that the agent should use the vanilla DQN algorithm without any modifications. Unlike Double DQN (DDQN), vanilla DQN uses the same network for both action selection and value estimation, which can lead to overestimation of Q-values. DDQN mitigates this issue by using the target network for value estimation while using the main network for action selection. The `--epsilon-decay` flag specifies the decay rate for the epsilon value, which controls the exploration-exploitation trade-off during training. The other hyperparameters are set to their default values.

The hyperparameters values are as follows:

- **Batch size:** 32

- **Discount factor:** 0.99
- **Learning rate:** 0.0001
- **Memory size:** 100,000
- **Epsilon start:** 1
- **Epsilon decay:** 0.99
- **Epsilon min:** 0.015
- **PER alpha:** 0
- **PER beta:** 1
- **N Steps:** 1
- **Update Period:** 1
- **Number of episodes:** 500
- **Seed:** 42
- **Evaluation episodes:** 10

The training process consists of 500 episodes, and the evaluation reward is calculated as the average reward over 10 evaluation episodes. The training process is performed using the Trainer class, which handles the training loop and updates the model.

3.1.2 Training Curves

The training curves for the task 1 are shown in Figure 1.



Figure 1: Training curves for the CartPole-v1 environment.

3.1.3 Testing Commands

We can test the trained agent using the following command:

```
1 python3 tester.py --env cartpole --exp report --model  
  ↪ ./results/cartpole/report/best_model.pt --episodes 30
```

This can achieve the score over than 480 in 30 consecutive episodes.

3.2 Task 2: ALE/Pong-v5 with vanilla DQN

3.2.1 Training Commands

For the task 2, which required to use the vanilla DQN algorithm, I used the following command to train the agent:

```
1 python3 trainer.py --env pong --exp report --vanilla --episodes 2500  
  ↪ --eval-episodes 2
```

The command specifies the environment as ALE/Pong-v5, the experiment name as `report`, and the `--vanilla` flag indicates that the agent should use the vanilla DQN algorithm without any modifications. Unlike Double DQN (DDQN), vanilla DQN uses the same network for both action selection and value estimation, which can lead to overestimation of Q-values. DDQN mitigates this issue by using the target network for value estimation while using the main network for action selection. The `--episodes` flag specifies the maximum number of episodes for training, which is set to 2500 in this case. The `--eval-episodes` flag specifies the number of evaluation episodes, which is set to 2 in this case. The other hyperparameters are set to their default values.

The hyperparameters values are as follows:

- **Batch size:** 32
- **Discount factor:** 0.99
- **Learning rate:** 0.000025
- **Memory size:** 100,000
- **Epsilon start:** 1
- **Epsilon decay:** 0.99999
- **Epsilon min:** 0.005
- **PER alpha:** 0
- **PER beta:** 1
- **N Steps:** 1
- **Update Period:** 1

- **Number of episodes:** 500
- **Seed:** 42
- **Evaluation episodes:** 2

3.2.2 Training Curves

The training curves for the task 2 are shown in Figure 2.



Figure 2: Training curves for the ALE/Pong-v5 environment.

We can see that the agent can achieve the score of 19 in some of the evaluation episodes.

3.2.3 Testing Commands

We can test the trained agent using the following command:

```
python3 tester.py --env pong --exp report --model
↪ ./results/pong/report/best_model.pt --episodes 20
```

This can achieve the score of 19.45 in 20 consecutive episodes.

3.3 Task 3: ALE/Pong-v5 with DQN Variants

3.3.1 Training Commands

For the task 3, which required to use the DQN variants.

However, I use the vanilla DQN algorithm for the ALE/Pong-v5 environment and **it can achieve the score of 21 (evaluated in 30 consecutive episodes) in about 140k environment steps.**

The training commands are as follows:

```
python3 trainer.py --env pong --exp task3 --vanilla --update-period 1
↪ --n-step 1
```

The command specifies the environment as ALE/Pong-v5, the experiment name as `task3`, and the `--vanilla` flag indicates that the agent should use the vanilla DQN algorithm without any modifications.

The other hyperparameters are set to their default values. The hyperparameters values are as follows:

- **Batch size:** 32
- **Discount factor:** 0.99
- **Learning rate:** 0.0001
- **Memory size:** 100,000
- **Epsilon start:** 1
- **Epsilon decay:** 0.99999
- **Epsilon min:** 0.005
- **PER alpha:** 0
- **PER beta:** 1
- **N Steps:** 1
- **Update Period:** 1
- **Number of episodes:** 500
- **Seed:** 42
- **Evaluation episodes:** 10

3.3.2 Training Curves

The training curves for the task 3 are shown in Figure 3.

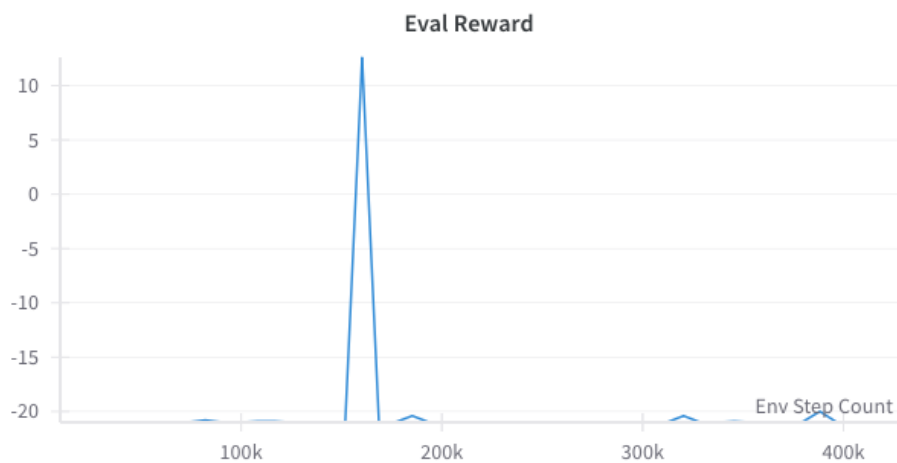


Figure 3: Training curves for the ALE/Pong-v5 environment.

We can see that the evaluation result shows that the agent can not achieve well performance in the ALE/Pong-v5 environment. It always get the score of -21 except a special snapshot that the agent can get the score over than 10. By the way, this evaluation score during the training process is the average score of 10 episodes.

After inspecting the special snapshot, I found that the agent can get the score of 21 in some episodes by moving to a specific position and hitting the ball at the right time without moving the paddle again. However, in the other episodes, the agent can not hit any ball and the score is -21 .

3.3.3 Testing Commands

Therefore, during the testing process, if we use this snapshot to test the agent in the ALE/Pong-v5 environment, the agent can get the score of 21 in 30 consecutive episodes by the following testing command:

```
python3 tester.py --env pong --exp task3 --model
↪ ./results/pong/task3/best_model.pt --seed 92439410 --episodes 30
```

The command specifies the environment as ALE/Pong-v5, the experiment name as task3, and the `--model` flag indicates that the agent should use the model saved in the `best_model.pt` file, you should change the path to your own path. The `--seed` flag specifies the first random seed for the environment, and the following seed for the environment will be `seed+episode_number`. That is to say, the first episode will use the seed 92439410, the second episode will use the seed 92439411, and so on. The `--episodes` flag specifies the number of evaluation episodes, which is set to 30 in this case.

3.4 Analyze each Technique

The following sections outline the key techniques employed in the DQN algorithm and their impact on the training process. Due to the extensive number of experiments conducted, including numerous subtle variations in hyperparameters, this report focuses on describing the techniques and analyzing their effects rather than presenting all training curves. For further details or to reproduce the results, please refer to the implementation in the `trainer.py` and `dqn.py` files.

3.4.1 Double DQN

Vanilla DQN uses a single target network to estimate the Q-value of the action it also selected, which tends to overestimate action values because both selection and evaluation share the same (noisy) numbers. Double DQN (DDQN) keeps the overall architecture and loss function the same but splits those two roles: the online network picks the action $a = \arg \max_a Q_{\text{online}}(s', a)$, while the target network supplies its value $Q_{\text{target}}(s', a)$ for the TD target. By decoupling selection from evaluation, DDQN dramatically reduces this positive bias, yielding more accurate value estimates, stabler learning curves, and usually better final performance—especially in environments with many actions or large reward variance—without adding extra networks or hyper-parameters beyond the standard target-network update already present in DQN.

3.5 Prioritized Experience Replay

Prioritized Experience Replay (PER) replaces the uniform-random sampling used in vanilla DQN’s replay buffer with a probability $P_i \propto |\delta_i|^\alpha$, where δ_i is the transition’s latest TD-error and $\alpha \in [0, 1]$ controls how aggressively "surprising" experiences are favored. Transitions that the network currently mis-predicts (large $|\delta|$) are therefore replayed more often, accelerating the correction of large errors and speeding convergence. Because this biased sampling breaks the i.i.d. assumption, PER attaches an importance-sampling weight $w_i = (\frac{1}{N} \frac{1}{P_i})^\beta$ (with annealed $\beta \in [0, 1]$) to each gradient to recover an unbiased estimate of the expected update. In practice, PER improves sample efficiency and final scores on many Atari and continuous-control tasks, costs only an extra log-time lookup via a SumTree or segment tree, and stacks well with other improvements like Double DQN or dueling networks—but it introduces two extra hyper-parameters (α, β) and can over-focus on a small subset of transitions if α is set too high.

3.6 Multi-Step Learning

Multi-step (n -step) learning replaces DQN’s single-step TD target with an n -step return that rolls rewards forward for n steps before bootstrapping:

$$G_t^{(n)} = r_t + \gamma r_{t+1} + \dots + \gamma^{n-1} r_{t+n-1} + \gamma^n \max_{a'} Q_{\text{target}}(s_{t+n}, a').$$

This amplifies the learning signal by injecting several real rewards before any bootstrapping noise, letting the agent propagate credit faster through time and making early training less sensitive to inaccurate value estimates. Because the return mixes Monte-Carlo (more accurate but high-variance) and TD (more biased but low-variance) signals, n controls a bias-variance trade-off: $n = 1$ reduces to vanilla DQN, $n \rightarrow \infty$ becomes a full Monte-Carlo target, and typical values $n \in [3, 10]$ give the best of both worlds. Multi-step targets are cheap to compute in a replay buffer (store cumulative discounted reward and terminal flag) and combine smoothly with other upgrades such as Double DQN and PER, often yielding faster convergence and higher final scores—especially in sparse-reward or long-horizon tasks—without adding new networks or extra hyper-parameters beyond the choice of n .

3.7 Additional Training Tricks

3.7.1 Multiple Evaluation Episodes

I found that the evaluation process is stochastic and the evaluation score is not stable during the training process. This means that the evaluation score can be affected by the random seed used in the environment and will cause the evaluation score to be unexpectedly high or low. This will cause the stored best model to be not the best model in the training process. Therefore, I used multiple evaluation episodes to calculate the average evaluation score during the training process. This can help to reduce the variance of the evaluation score and make the evaluation score more stable. This can be achieved by using the `-eval-episodes` flag in the training command. The `-eval-episodes` flag specifies the number of evaluation episodes, which is set to 10 by default.

3.7.2 Learning Rate Scheduler

I found that the learning rate is a very important hyper-parameter in the training process. The learning rate controls the step size of the gradient descent algorithm and can affect the

convergence speed and stability of the training process. From my experiments, I found that if the learning rate is too high, the training process will stuck at a suboptimal solution and the evaluation score will not improve no matter how many episodes are trained. On the other hand, if the learning rate is too low, the training process will be very slow and the evaluation score, but the evaluation score will keeps improving. Therefore, I used a learning rate scheduler to reduce the learning rate during the training process, which can help to improve the convergence speed and stability of the training process.

References

- [1] H. van Hasselt, A. Guez, and D. Silver, “Deep reinforcement learning with double q-learning,” in *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence (AAAI-16)*, 2016, pp. 2094–2100. [Online]. Available: <https://arxiv.org/abs/1509.06461>.
- [2] T. Schaul, J. Quan, I. Antonoglou, and D. Silver, “Prioritized experience replay,” in *4th International Conference on Learning Representations (ICLR)*, arXiv:1511.05952, 2016. [Online]. Available: <https://arxiv.org/abs/1511.05952>.
- [3] B. Daley, M. White, and M. C. Machado, “Averaging n -step returns reduces variance in reinforcement learning,” in *Proceedings of the 41st International Conference on Machine Learning (ICML)*, arXiv:2402.03903, 2024. [Online]. Available: <https://arxiv.org/abs/2402.03903>.